

A surrogate's gradients agree with the solver only inside the identifiable subspace
mean cosine inside +0.504, outside -0.001 (outside is flat across a 33× budget, so it is not undertraining)

